

Document Number: DXXXXR0
Date: 2019-03-23
Reply to: Daniel Sunderland
Sandia National Laboratories
dsunder@sandia.gov

Alternate `atomic_ref` for Non-Lockfree Types

```
git clone git@github.com:dsunder/draft.git dsunder-draft
cd dsunder-draft
git checkout atomic_ref_alt
```

Motivation

Enable implementations of `atomic_ref` which do not require a global lock array while perserving current behavior as much as possible.

Proposed Wording

This font is used to provide guidance to the editors.

Make the following changes in [atomics.ref.generic].

```
namespace std {
    template<class T, class Lock = mutex> struct atomic_ref {
    private:
        T* ptr;           // exposition only
        Lock* lock; // exposition only
    public:
        using value_type = T;
        static constexpr bool is_always_lock_free = implementation-defined;
        static constexpr bool never_requires_user_lock = implementation-defined;
        static constexpr size_t required_alignment = implementation-defined;

        static bool is_lock_free(const T &) noexcept;
        static bool requires_user_lock(const T &) noexcept;

        atomic_ref& operator=(const atomic_ref&) = delete;

        explicit atomic_ref(T&, Lock* = nullptr);
        atomic_ref(const atomic_ref&) noexcept;

        T operator=(T) const noexcept;
        operator T() const noexcept;

        bool is_lock_free() const noexcept;
        void store(T, memory_order = memory_order_seq_cst) const noexcept;
        T load(memory_order = memory_order_seq_cst) const noexcept;
        T exchange(T, memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_weak(T&, T,
                                   memory_order, memory_order) const noexcept;
        bool compare_exchange_strong(T&, T,
                                     memory_order, memory_order) const noexcept;
        bool compare_exchange_weak(T&, T,
                                   memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_strong(T&, T,
                                     memory_order = memory_order_seq_cst) const noexcept;
    };
}
```

Make the following changes in [atomics.ref.operations].

```
static constexpr bool is_always_lock_free;
```

- ¹ The static data member `is_always_lock_free` is true if the `atomic_ref` type's operations are always lock-free on objects aligned to required alignment, and false otherwise.

```
static constexpr bool never_requires_user_lock;
```

- ² The static data member `never_requires_user_lock` is true if the `atomic_ref` type's operations never require the user to provide a valid pointer to a `Lock` object.

```
static bool is_lock_free(T& obj); noexcept
```

Returns: Returns true if atomic operations on the object referenced by `obj` can be lock-free.

```
static requires_user_lock(T& obj); noexcept
```

Returns: Returns true if atomic_ref requires the user to provide a valid pointer to a Lock object when constructing an atomic_ref from obj. [Note: Implementations could provide ways for atomic_ref to use implementation defined locking mechanisms when is_lock_free(obj) is false. — end note]

```
atomic_ref(T& obj\added{, Lock* lk = nullptr});
```

3 ~~*Requires:* The referenced object shall be aligned to required_alignment~~ *Expects:* The type Lock meets the Cpp17BasicLockable requirements and lk points to the the same address for all atomic_ref which reference the object referenced by obj. If is_lock_free(obj) is false and requires_user_lock(obj) is false then lk points to either a valid Lock object or nullptr. If requires_user_lock is true then lk points to valid Lock object.

4 *Effects:* ~~Constructs an atomic reference that references the object.~~
Equivalent to:

```
ptr = &obj;
lock = is_lock_free(obj) ? nullptr : lk;
```

5 *Throws:* Nothing.

```
atomic_ref(const atomic_ref& ref) noexcept;
```

6 *Effects:* ~~Constructs an atomic reference that references the object referenced by ref.~~ Equivalent to:

```
ptr = ref.obj;
lock = ref.lock;
```

```
void store(T desired, memory_order order = memory_order_seq_cst) const noexcept;
```

7 ~~*Requires:*~~ *Expects:* The order argument shall not be memory_order_consume, memory_order_acquire, nor memory_order_acq_rel.

8 ~~Atomically replaces the value referenced by *ptr with the value of desired.~~ *Effects:* If lock is equivalent to nullptr then the operation is equivalent to atomically performing the following:

```
memcpy(ptr, &desired, sizeof(T));
```

Otherwise, equivalent to:

```
lock->lock();
memcpy(ptr, &desired, sizeof(T));
lock->unlock();
```

Memory is affected according to the value of order.

```
T load(memory_order order = memory_order_seq_cst) const noexcept;
```

9 ~~*Requires:*~~ *Expects:* The order argument shall not be memory_order_release nor memory_order_acq_rel.

10 *Effects:* If lock is equivalent to nullptr then the operation is equivalent to atomically performing the following:

```
T result;
memcpy(&result, ptr, sizeof(T));
return result;
```

Otherwise, equivalent to:

```
T result;
lock->lock();
memcpy(&result, ptr, sizeof(T));
lock->unlock();
return result;
```

Memory is affected according to the value of order.

11 ~~*Returns:* Atomically returns the value referenced by *ptr.~~

```
T exchange(T desired, memory_order order = memory_order_seq_cst) const noexcept;
```

12 ~~Atomically replaces the value referenced by *ptr with desired.~~ *Effects:* If lock is equivalent to nullptr then the operation is equivalent to atomically performing the following:

```

    T result;
    memcpy(&result, ptr, sizeof(T));
    memcpy(ptr, &desired, sizeof(T));
    return result;

```

Otherwise, equivalent to:

```

    T result;
    lock->lock();
    memcpy(&result, ptr, sizeof(T));
    memcpy(ptr, &desired, sizeof(T));
    lock->unlock();
    return result;

```

Memory is affected according to the value of `order`. This operation is an atomic read-modify-write operation (??).

13 *Returns:* Atomically returns the value referenced by `*ptr` immediately before the effects.

```

bool compare_exchange_weak(T& expected, T desired,
                           memory_order success, memory_order failure) const noexcept;

bool compare_exchange_strong(T& expected, T desired,
                             memory_order success, memory_order failure) const noexcept;

bool compare_exchange_weak(T& expected, T desired,
                           memory_order order = memory_order_seq_cst) const noexcept;

bool compare_exchange_strong(T& expected, T desired,
                             memory_order order = memory_order_seq_cst) const noexcept;

```

14 ~~*Requires:*~~ *Expects:* The failure argument shall not be `memory_order_release` nor `memory_order_acq_rel`.

15 *Effects:* When only one `memory_order` argument is supplied, the value of `success` is `order`, and the value of `failure` is `order` except that a value of `memory_order_acq_rel` shall be replaced by the value `memory_order_acquire` and a value of `memory_order_release` shall be replaced by the value `memory_order_relaxed`.

If `lock` is equivalent to `nullptr` then the operation is equivalent to atomically performing the following:

```

    T old;
    memcpy(&old, ptr, sizeof(T));
    bool result = 0 == memcmp(&expected, &old, sizeof(T));
    if (result) memcpy(ptr, &desired, sizeof(T));
    else memcpy(&expected, &old, sizeof(T));
    return result;

```

Otherwise, equivalent to:

```

    T old;
    lock->lock();
    memcpy(&old, ptr, sizeof(T));
    bool result = 0 == memcmp(&expected, &old, sizeof(T));
    if (result) memcpy(ptr, &desired, sizeof(T));
    else memcpy(&expected, &old, sizeof(T));
    lock->unlock();
    return result;

```

16 Let `R` be the return value of the operation. If and only if `R` is `true`, memory is affected according to the value of `success`, and if `R` is `false`, memory is affected according to the value of `failure`.

Retrieves the value in `expected`. It then atomically compares the value representation of the value referenced by `*ptr` for equality with that previously retrieved from `expected`, and if `true`, replaces the value referenced by `*ptr` with that in `desired`. When only one `memory_order` argument is supplied, the value of `success` is `order`, and the value of `failure` is `order` except that a value of `memory_order_acq_rel` shall be replaced by the value `memory_order_acquire` and a value of `memory_order_release` shall be replaced by the value `memory_order_relaxed`. If and only if the comparison is `false` then, after the atomic operation, the value in `expected` is replaced by the value read from the value referenced

by **ptr* during the atomic comparison. If ~~the operation returns~~*R* is true, these operations are atomic read-modify-write operations (??) on the value referenced by **ptr*. Otherwise, these operations are atomic load operations on that memory.

17 *Returns:* The result of the comparison.

18 *Remarks:* A weak compare-and-exchange operation may fail spuriously. That is, even when the contents of memory referred to by *expected* and *ptr* are equal, it may return false and store back to *expected* the same memory contents that were originally there. [Note: This spurious failure enables implementation of compare-and-exchange on a broader class of machines, e.g., load-locked store-conditional machines. A consequence of spurious failure is that nearly all uses of weak compare-and-exchange will be in a loop. When a compare-and-exchange is in a loop, the weak version will yield better performance on some platforms. When a weak compare-and-exchange would require a loop and a strong one would not, the strong one is preferable. — end note]

Make the following changes in [atomics.ref.int].

```
namespace std {
    template<Lock> struct atomic_ref<integral, Lock> {
    private:
        integral* ptr;           // exposition only
        Lock* lock; // exposition only
    public:
        using value_type = integral;
        using difference_type = value_type;
        static constexpr bool is_always_lock_free = implementation-defined;
        static constexpr bool never_requires_user_lock = implementation-defined;
        static constexpr size_t required_alignment = implementation-defined;

        static bool is_lock_free(const T &) noexcept;
        static bool requires_user_lock(const T &) noexcept;

        atomic_ref& operator=(const atomic_ref&) = delete;

        explicit atomic_ref(integral&, Lock* = nullptr);
        atomic_ref(const atomic_ref&) noexcept;

        integral operator=(integral) const noexcept;
        operator integral() const noexcept;

        bool is_lock_free() const noexcept;
        void store(integral, memory_order = memory_order_seq_cst) const noexcept;
        integral load(memory_order = memory_order_seq_cst) const noexcept;
        integral exchange(integral,
                           memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_weak(integral&, integral,
                                   memory_order, memory_order) const noexcept;
        bool compare_exchange_strong(integral&, integral,
                                      memory_order, memory_order) const noexcept;
        bool compare_exchange_weak(integral&, integral,
                                    memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_strong(integral&, integral,
                                      memory_order = memory_order_seq_cst) const noexcept;

        integral fetch_add(integral,
                           memory_order = memory_order_seq_cst) const noexcept;
        integral fetch_sub(integral,
                           memory_order = memory_order_seq_cst) const noexcept;
        integral fetch_and(integral,
                           memory_order = memory_order_seq_cst) const noexcept;
```

```

    integral fetch_or(integral,
                      memory_order = memory_order_seq_cst) const noexcept;
    integral fetch_xor(integral,
                      memory_order = memory_order_seq_cst) const noexcept;

    integral operator++(int) const noexcept;
    integral operator--(int) const noexcept;
    integral operator++() const noexcept;
    integral operator--() const noexcept;
    integral operator+=(integral) const noexcept;
    integral operator-=(integral) const noexcept;
    integral operator&=(integral) const noexcept;
    integral operator|=(integral) const noexcept;
    integral operator^=(integral) const noexcept;
};
}

```

Make the following changes in *[atomics.ref.float]*.

```

namespace std {
    template<Lock> struct atomic_ref<floating-point, Lock> {
    private:
        floating-point* ptr; // exposition only
        Lock* lock; // exposition only
    public:
        using value_type = floating-point;
        using difference_type = value_type;
        static constexpr bool is_always_lock_free = implementation-defined;
        static constexpr bool never_requires_user_lock = implementation-defined;
        static constexpr size_t required_alignment = implementation-defined;

        static bool is_lock_free(const T &) noexcept;
        static bool requires_user_lock(const T &) noexcept;

        atomic_ref& operator=(const atomic_ref&) = delete;

        explicit atomic_ref(floating-point&, Lock* = nullptr);
        atomic_ref(const atomic_ref&) noexcept;

        floating-point operator=(floating-point) noexcept;
        operator floating-point() const noexcept;

        bool is_lock_free() const noexcept;
        void store(floating-point, memory_order = memory_order_seq_cst) const noexcept;
        floating-point load(memory_order = memory_order_seq_cst) const noexcept;
        floating-point exchange(floating-point,
                               memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_weak(floating-point&, floating-point,
                                   memory_order, memory_order) const noexcept;
        bool compare_exchange_strong(floating-point&, floating-point,
                                      memory_order, memory_order) const noexcept;
        bool compare_exchange_weak(floating-point&, floating-point,
                                    memory_order = memory_order_seq_cst) const noexcept;
        bool compare_exchange_strong(floating-point&, floating-point,
                                      memory_order = memory_order_seq_cst) const noexcept;

        floating-point fetch_add(floating-point,
                                memory_order = memory_order_seq_cst) const noexcept;
        floating-point fetch_sub(floating-point,
                                memory_order = memory_order_seq_cst) const noexcept;

        floating-point operator+=(floating-point) const noexcept;
        floating-point operator-=(floating-point) const noexcept;
    };
}

```

```
};
}
```

Make the following changes in [atomics.ref.pointer].

```
namespace std {
template<class T, Lock> struct atomic_ref<T*, Lock> {
private:
    T** ptr; // exposition only
    Lock* lock; // exposition only
public:
    using value_type = T*;
    using difference_type = ptrdiff_t;
    static constexpr bool is_always_lock_free = implementation-defined;
    static constexpr bool never_requires_user_lock = implementation-defined;
    static constexpr size_t required_alignment = implementation-defined;

    static bool is_lock_free(const T &) noexcept;
    static bool requires_user_lock(const T &) noexcept;

    atomic_ref& operator=(const atomic_ref&) = delete;

    explicit atomic_ref(T*&, Lock* = nullptr);
    atomic_ref(const atomic_ref&) noexcept;

    T* operator=(T*) const noexcept;
    operator T*() const noexcept;

    bool is_lock_free() const noexcept;
    void store(T*, memory_order = memory_order_seq_cst) const noexcept;
    T* load(memory_order = memory_order_seq_cst) const noexcept;
    T* exchange(T*, memory_order = memory_order_seq_cst) const noexcept;
    bool compare_exchange_weak(T*&, T*,
                               memory_order, memory_order) const noexcept;
    bool compare_exchange_strong(T*&, T*,
                                 memory_order, memory_order) const noexcept;
    bool compare_exchange_weak(T*&, T*,
                               memory_order = memory_order_seq_cst) const noexcept;
    bool compare_exchange_strong(T*&, T*,
                                 memory_order = memory_order_seq_cst) const noexcept;

    T* fetch_add(difference_type, memory_order = memory_order_seq_cst) const noexcept;
    T* fetch_sub(difference_type, memory_order = memory_order_seq_cst) const noexcept;

    T* operator++(int) const noexcept;
    T* operator--(int) const noexcept;
    T* operator++() const noexcept;
    T* operator--() const noexcept;
    T* operator+=(difference_type) const noexcept;
    T* operator-=(difference_type) const noexcept;
};
}
```